Package 'snht'

November 3, 2014

Type Package

Title Standard Normal Homogeneity Test

Version 1.0

Date 2014-09-23

Imports MASS, zoo, plyr, reshape

Author Josh Browning

Maintainer Josh Browning <jbrownin@mines.edu>

Description Robust and non-robust SNHT tests for changepoint detection.

License GPL-3

R topics documented:

`snht-package` *Robust and Non-Robust Standard Normal Homogeneity Test*

Description

Computes test statistics for the SNHT and robust SNHT test. Additionally, users may supply a custom function for estimating the mean and standard deviation, and this function will be used for computing the test statistic.

Details

Package:	snht
Type:	Package
Version:	1.0
Date:	2014-09-23
License:	What license is it under?

The main function is snht, which then calls the other functions in this package. However, users may also wish to call robustSNHT which allows for a custom estimator function.

Author(s)

Josh Browning

Maintainer: Josh Browning <jbrownin@mines.edu>

References

L. Haimberger. Homogenization of radiosonde temperature time series using innovation statistics. Journal of Climate, 20(7): 1377-1403, 2007.

`pairwiseSNHT` *Pairwise Standard Normal Homogeneity Test*

Description

This function performs a pairwise standard normal homogeneity test on the data supplied, as described in Menne & Williams (2009).

Usage

```
pairwiseSNHT(data, dist, k, period, crit=100, returnStat=FALSE, ...)
```

Arguments

`data`	The data to be analyzed for changepoints. It must be a data.frame and contain either two or three columns. The mandatory columns are data and location, named as such. The option column is time, and this argument will be passed to snht.
`dist`	A distance matrix which provides the distance between location i and location j. Rows and columns must be named with the locations in data. Note that non-symmetric distances may be used. In that case, neighbors for station i will be determined by the smallest values in the row of dist corresponding to i.
`k`	How many of the nearest neighbors should be used to construct pairwise difference time series? Note that more than k neighbors may be used if there are ties in the distances between locations.

period	The SNHT works by calculating the mean of the data on the previous period observations and the following period observations. Thus, this argument controls the window size for the test statistics.
crit	The critical value such that if the snht statistic is larger than crit, a changepoint is assumed to have occured. Defaults to 100, as recommended in Haimberger (see references).
returnStat	See return value. If TRUE, the snht statistics for each time point and for each difference pair are returned.
...	Additional arguments to pass to the snht function (such as robust, time, or estimator).

Details

The pairwise snht works with a set of time series. For each time series, it's closest k neighbors are determined, and a time series of the difference between each of those time series is created. The snht is then applied to each of these difference time series. Changepoints in one time series can be detected by searching for large values of the test statistic across all difference time series for a particular location.

The usefulness of the pairwise snht is that it removes any patterns in the data that could affect the basic snht. For example, seasonal and linear trends that exist globally will be removed from the difference series, and thus changepoints are more easily detected.

Value

If returnStat is TRUE, the snht statistics for each time point and for each difference pair are returned.

Otherwise, a named list is returned. The first element, data, contains the homogenized data in the same format as the supplied data. The second element, breaks, contains a data.frame where the first column is the location where a break occured, the second column is the time of the break, and the third column is the amount that data after the break was shifted by.

Author(s)

Josh Browning (jbrownin@mines.edu)

References

L. Haimberger. Homogenization of radiosonde temperature time series using innovation statistics. Journal of Climate, 20(7): 1377-1403, 2007.

Menne, M. J., & Williams Jr, C. N. (2009). Homogenization of temperature series via pairwise comparisons. Journal of Climate, 22(7), 1700-1717.

snht *Standard Normal Homogeneity Test*

Description

This function performs a standard normal homogeneity test on the data supplied. This test searches the data for potential changepoints.

Usage

```
snht(data, period, robust = F, time = NULL, ...)
robustSNHT(data, period, estimator = function(x, minObs = 5) {
    x = x[!is.na(x)]
    if (length(x) < minObs)
        return(c(NA, NA))
    if (max(table(x)) > length(x)/2)
        return(c(NA, NA))
    fit = MASS::huber(x)
    return(c(fit[[1]], fit[[2]]))
})
robustSNHTunequal(data, period, time, estimator = NULL)
```

Arguments

data	The data to be analyzed for changepoints.
period	The SNHT works by calculating the mean of the data on the previous period observations and the following period observations. Thus, this argument controls the window size for the test statistics.
robust	Flag indicating whether or not robust estimators should be used. If T, then Huber's robust estimator for the mean and variance will be used (see ?MASS::huber).
time	Numeric vector specifying times for the observations. If not supplied, it is assumed that each observation occurs on one time period. If supplied, then the algorithm will create a new dataset with the same number of observations for each time unit by adding missing values.
estimator	A custom function may be supplied to this function which computes estimates for the mean and standard deviation. The function should only take one argument (a numeric vector of data) and should return a vector of length two: the estimated center and spread. The huber function from MASS is implemented for the robust SNHT by default (along with some data quality checks).
...	Additional arguments to pass to the robustSNHT function.

Details

The SNHT works by calculating the mean of the data on the previous period and on the following period. The test statistic at each observation is then computed as described in Haimberger (2007).

Essentially, though, it just compares the means of these two periods and normalizes by the standard deviation.

Note: if there are not enough observations both before and after the current observation, no test is performed.

Large values of the test statistic suggests the presence of a changepoint. Haimberger (see references) suggests values larger than 100 should be considered changepoints.

Value

Returns a data.frame, with columns score, leftMean, and rightMean. Statistic is the SNHT test statistic described above, and leftMean (rightMean) are the means to the left (right) of the current observation.

Additionally, if time is supplied, then time is returned on the output data.frame. Note that new (missing) observations were introduced to the dataset to ensure the same number of observations occur per day.

Author(s)

Josh Browning (jbrownin@mines.edu)

References

L. Haimberger. Homogenization of radiosonde temperature time series using innovation statistics. Journal of Climate, 20(7): 1377-1403, 2007.

See Also

huber

Examples

```
data = rnorm(1000)
brk = sample(1000, size=1)
data[1:brk] = data[1:brk]-2
out = snht( data, period=50, robust=FALSE )
summary(out)

data = rnorm(1000)
time = 1:1000 + rnorm(1000)
brk = sample(1000, size=1)
data[1:brk] = data[1:brk]-2
out = snht( data, period=50, time=time, robust=FALSE )
summary(out)
```

Index